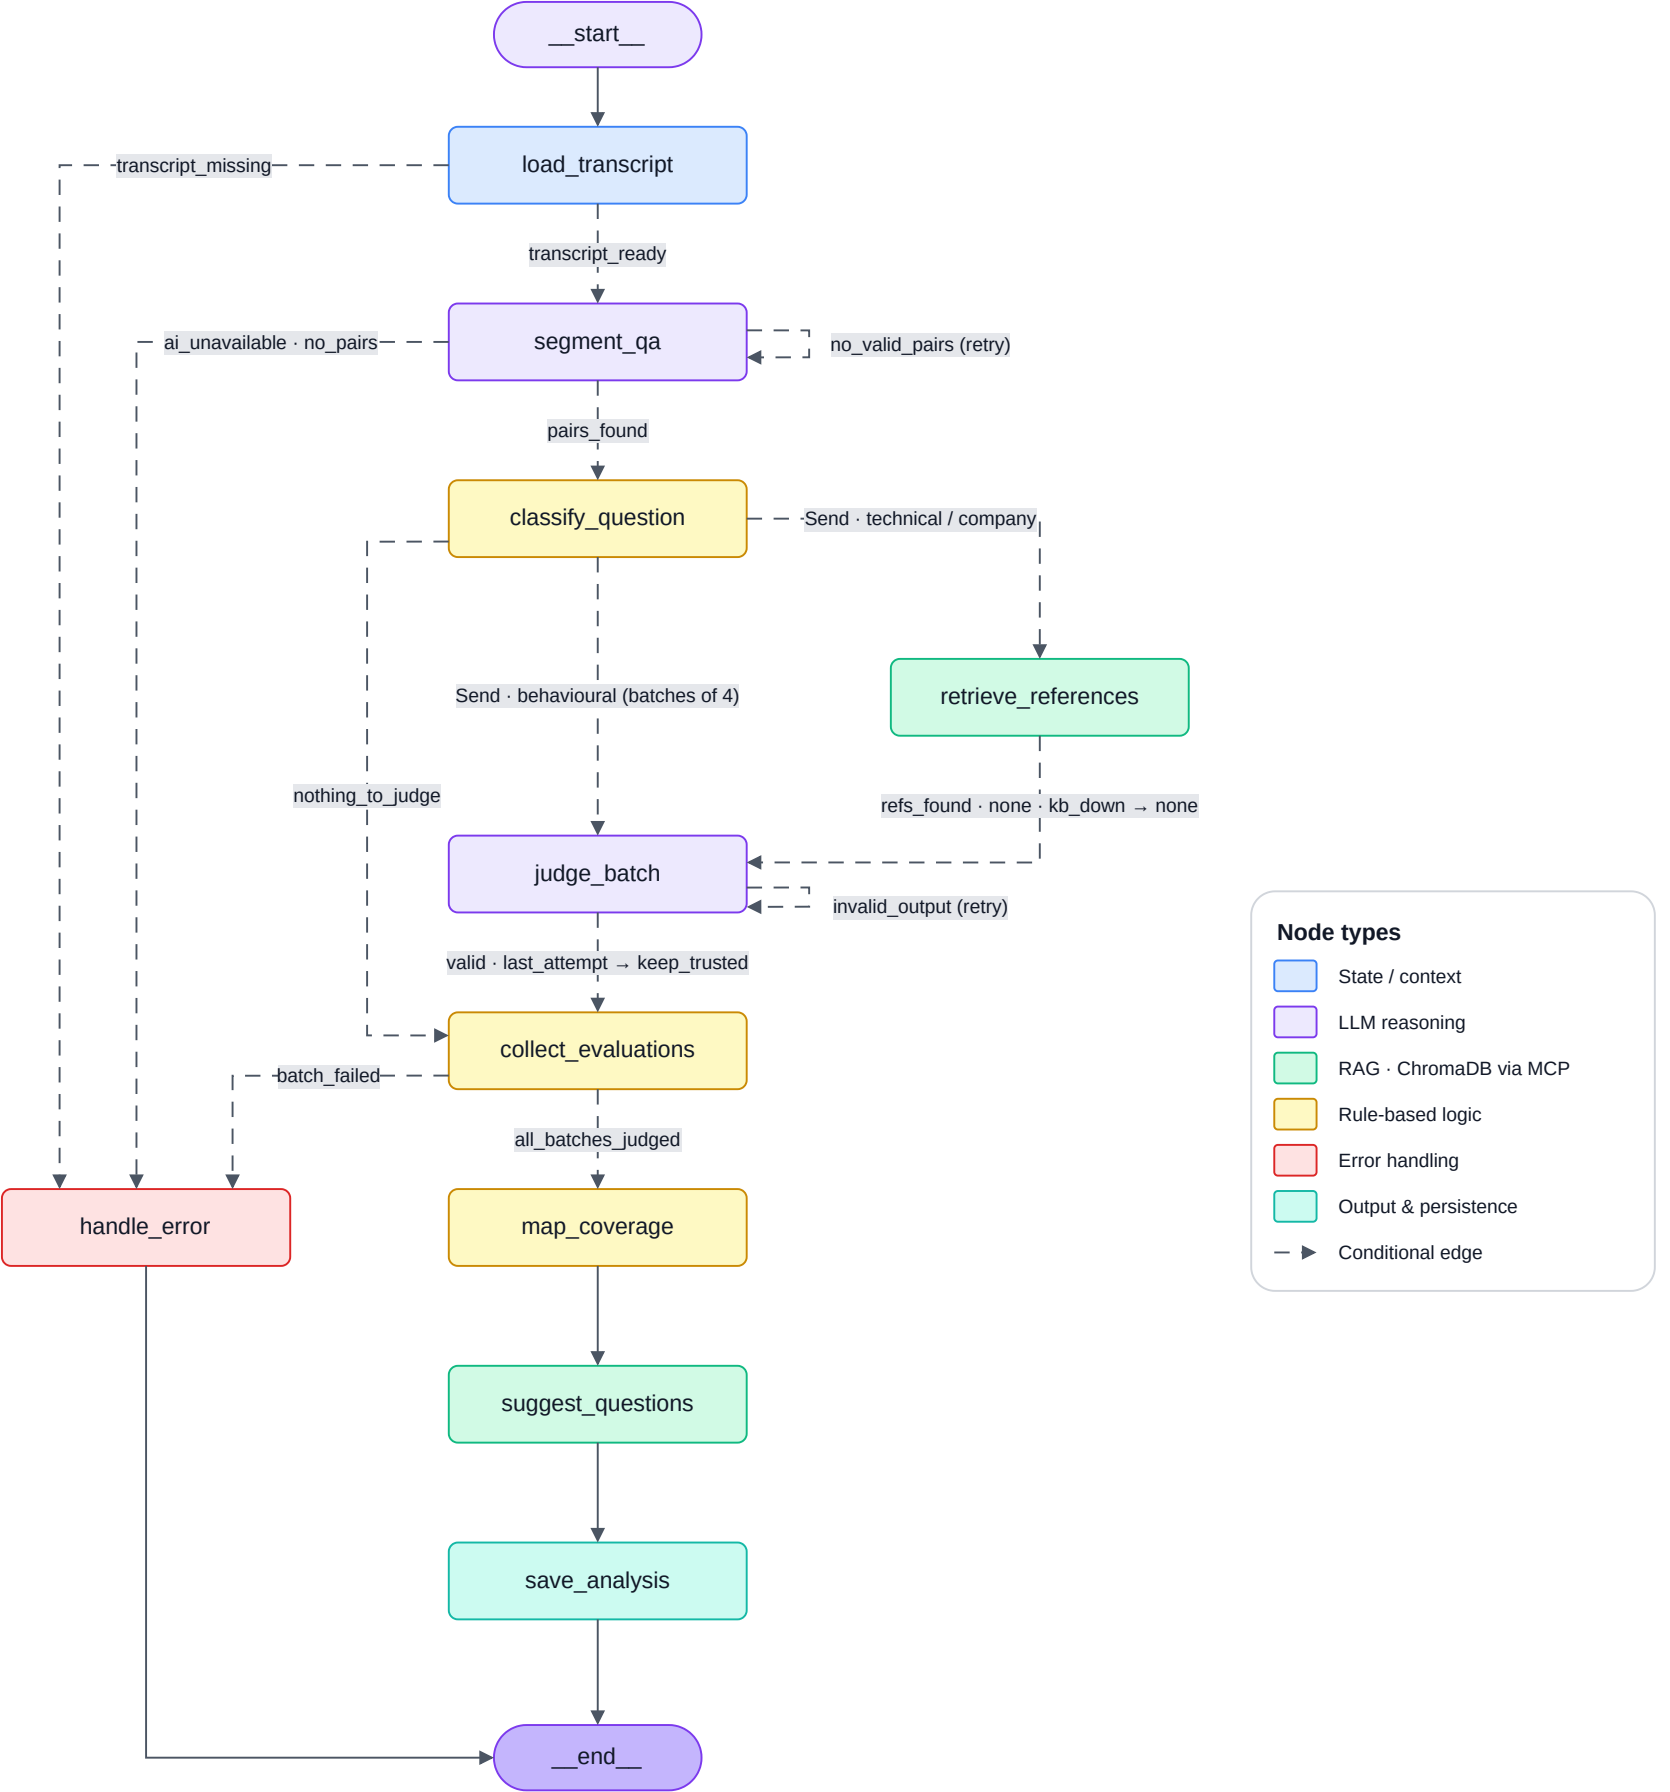
__start__
load_transcript
transcript_missing
transcript_ready
segment_qa
ai_unavailable · no_pairs
no_valid_pairs (retry)
pairs_found
classify_question
Send · technical / company
retrieve_references
Send · behavioural (batches of 4)
nothing_to_judge
refs_found · none · kb_down → none
judge_batch
invalid_output (retry)
valid · last_attempt → keep_trusted
collect_evaluations
batch_failed
all_batches_judged
handle_error
map_coverage
suggest_questions
save_analysis
__end__
Node types
State / context
LLM reasoning
RAG · ChromaDB via MCP
Rule-based logic
Error handling
Output & persistence
Conditional edge